

The NTNIX HASH Device Library

Last revision: August 2018

Expected completion time: under 60 mins

Introduction

You have been tasked to write a Linux library which will enable applications to use a fictional MD5 hashing device. This is managed by a kernel driver which provides a userspace interface via a character device. Your library will interact solely with this interface and no kernel programming is required. Applications using your library will be able to compute the digest of arbitrarily-sized buffers very efficiently, given the bulk of the work will be offloaded to the device instead of being performed by a CPU.

Library API

The API of this library consists of three functions (which you are expected to implement):

```
ntnx_hash_t *ntnx_hash_setup(void);  
char *ntnx_hash_compute(ntnx_hash_t *ctx, void *buf, size_t len);  
int ntnx_hash_destroy(ntnx_hash_t *ctx);
```

To use the library, an application must first create a context by calling `ntnx_hash_setup()`. Using the created context, it can compute MD5 digests by calling `ntnx_hash_compute()`. Apart from the context, `ntnx_hash_compute()` takes a `buf` of size `len` which must contain the data to be hashed. Finally, the application can destroy the context with `ntnx_hash_destroy()`. Multi-threaded applications should be able to share a context across threads

Return values

`ntnx_hash_setup()` returns a pointer to a `ntnx_hash_t` context. This should be opaque to the user and you may define it internally in your library as you see fit. In the case of errors, the function must return `NULL` and set `errno` appropriately.

`ntnx_hash_compute()` returns a pointer to a null-terminated array of 33 chars (including `NULL`) which will be allocated in the heap (by your library). It must be later `free()`d by the caller. Multi-threaded programs should be able to call this concurrently. Upon successful computation, the array must contain the 32-byte string representation of the MD5 digest for the buffer provided. In the case of errors, the function must return `NULL` and set `errno` appropriately.

`ntnx_hash_destroy()` should release all resources associated with the context. It returns zero on success or -1 on error, in which case it should set `errno` appropriately.

The `errno` codes to be used are left at your discretion, as well as the state of the system in the case of errors.

Character Device

Systems with a valid device and a correctly loaded kernel module will provide a character device on `/dev/ntnx_hash`. Your library must open this device upon context creation. On context destruction, it must close the device. While a context is open, it can issue the following `ioctl()`s to operate the device driver:

1) API version retrieval

```
#define NTNX_HASH_GET_API_VERSION 0
```

This takes the address of an unsigned int as an "out" parameter. The driver will fill it with its API version. For the purposes of this exercise, only version 1 exists. The `ioctl()` itself will return 0 on success and -1 on error. Your library must check it is compatible with the device driver API.

2) Hash computation

```
#define NTNX_HASH_COMPUTE 1
```

This takes a `struct ntnx_hash_compute` as defined below.

```
struct ntnx_hash_compute {  
    void *buf;        // pointer to the area for hashing  
    size_t len;        // length of area for checksumming  
    void *hash;        // pointer to the area for the computed hash  
};
```

The driver will offload the area of size `len` pointed to by `buf` to the device. The device will calculate the MD5 hash of the area and write the 32-byte string representation of the digest plus a `NUL` to the address pointed to by `hash`. The `ioctl()` itself will return 0 on success and -1 on error. The device driver expects invocations of this `ioctl()` to be serialised within a context.

Expected solution

Your solution should consist of two files:

1) `ntnx_hash.h`

This file should contain the headers for using your library.

2) `ntnx_hash.c`

This file should contain the implementation of your library.